Translating SQL to pandas. And back.

PyData NYC
November 24, 2014

Greg Reda

- @gjreda
- gregreda.com
- Studied economics
- Led data at GrubHub
- Data Scientist at Datascope Analytics

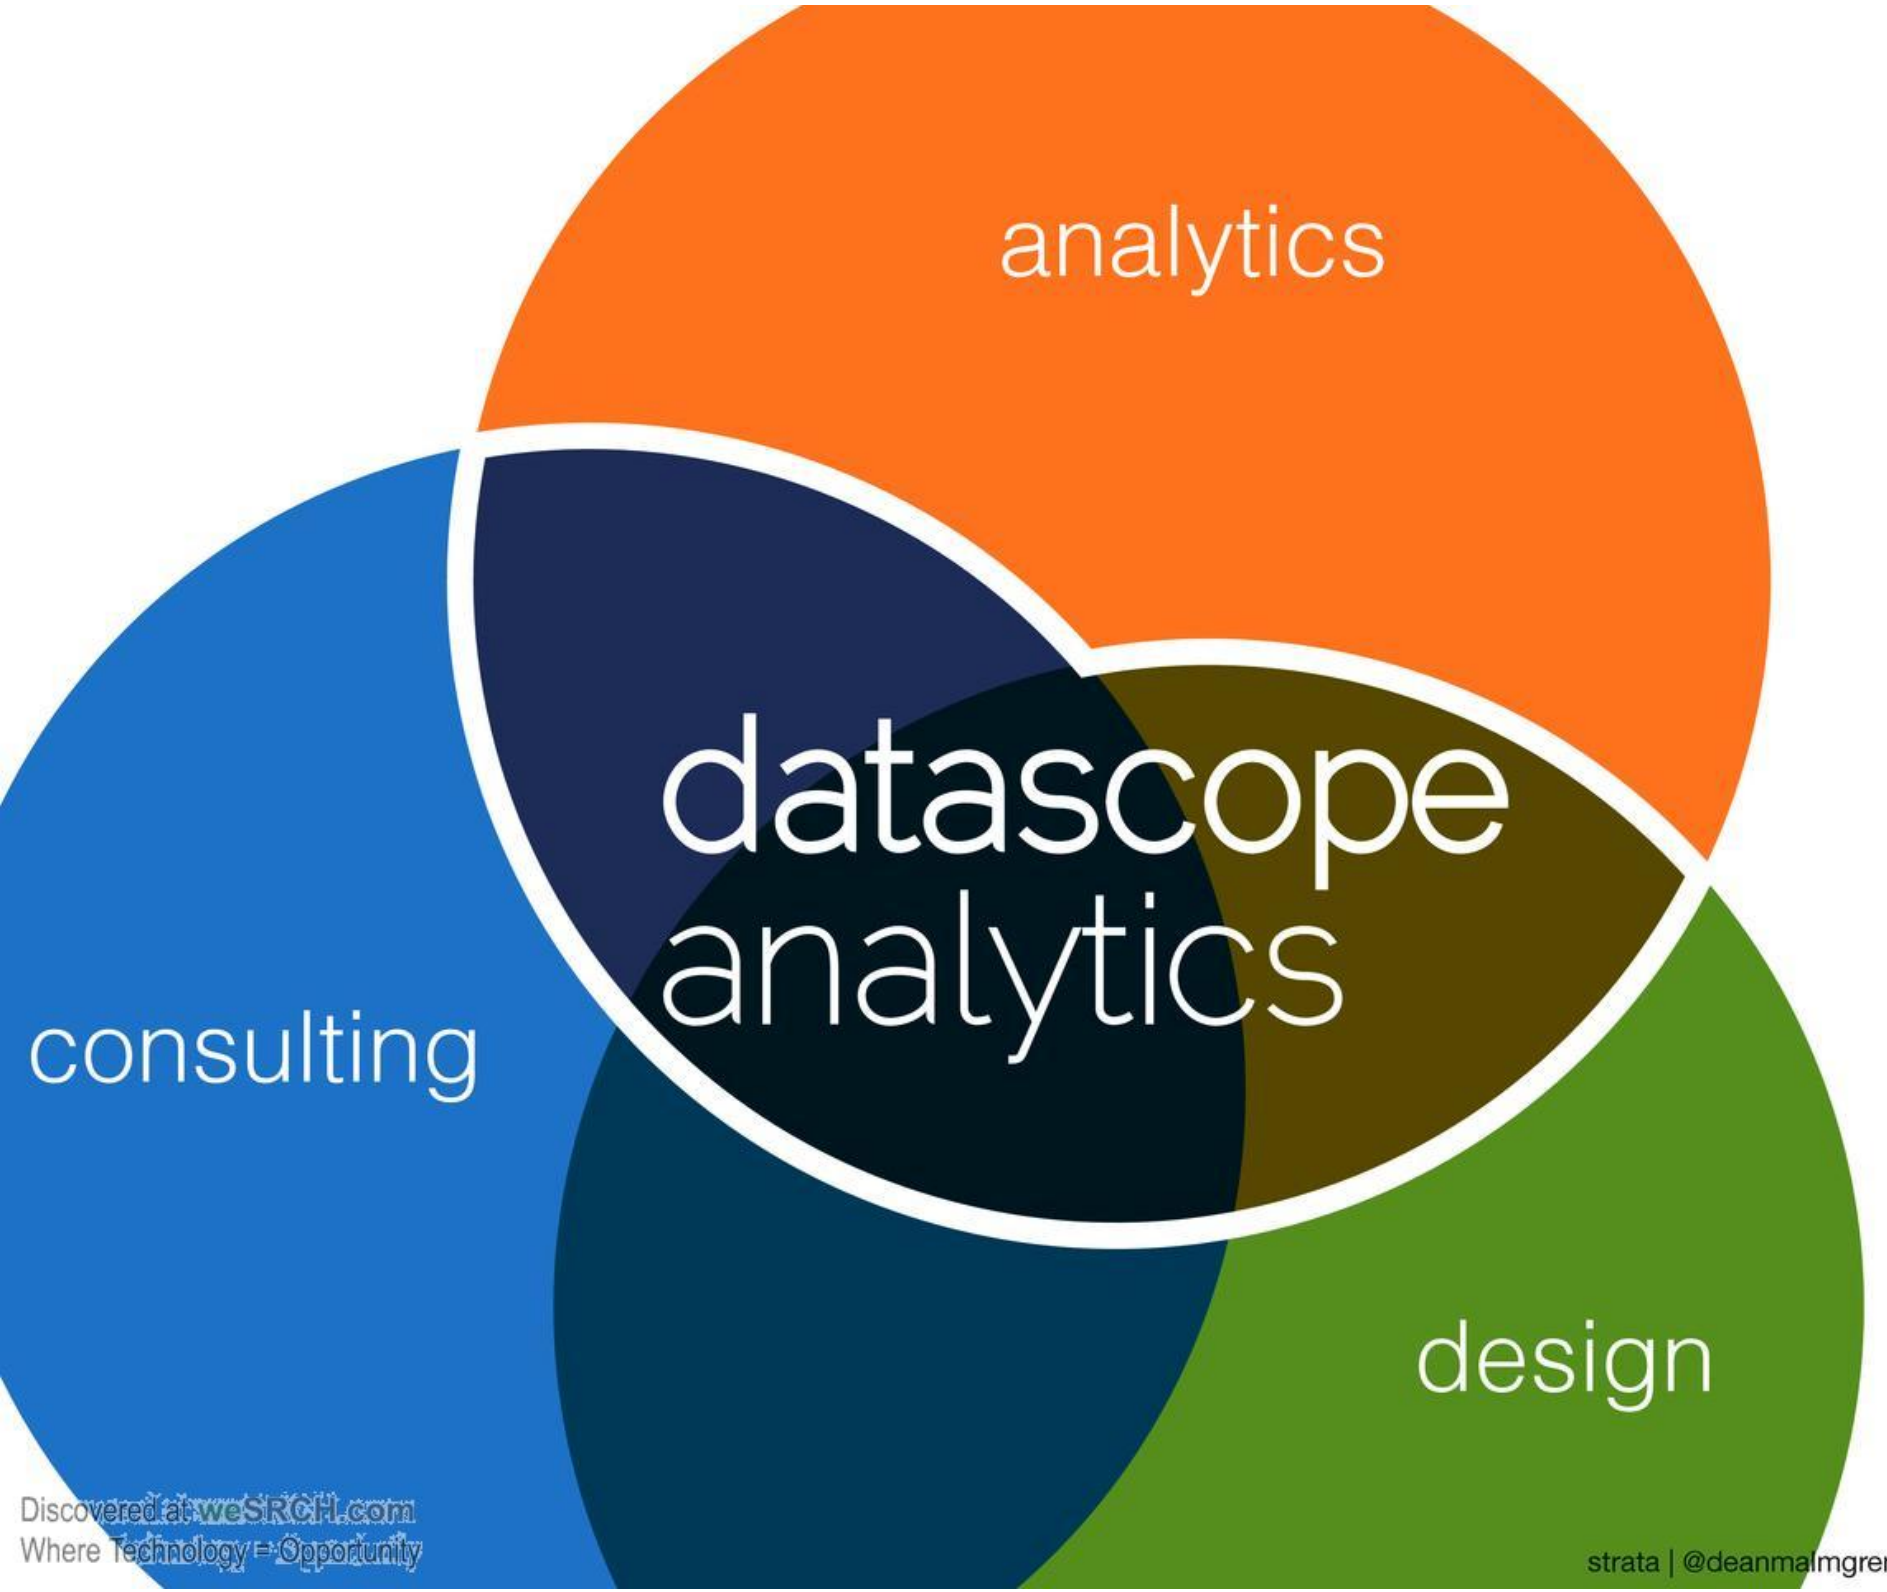

strata | @deanmalmgren | 2013 february

pandas

- Started by Wes McKinney in 2008
- Python lacked data *analysis* capabilities
- Built on top of NumPy (that means it's fast)
- 300+ contributors
- Big and active community led by Jeff Reback

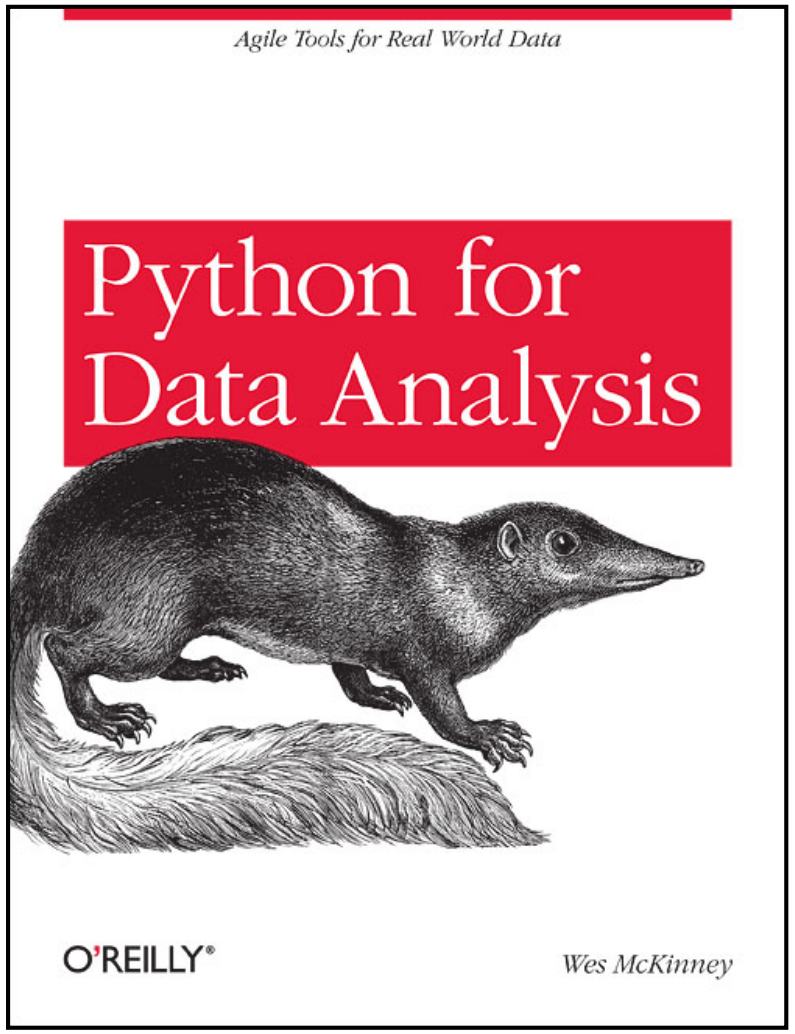


PANDA
ZOO
GIANT PANDA

pandas is PyData glue

What ***isn't*** pandas?

What we'll cover

- Data structures
  - Series & DataFrames
  - Indexes
- I/O - getting your data in and out of pandas
- Working with DataFrames
- Applied analysis (using IPython Notebook)

Series

In [2]:
```
# create a Series with an arbitrary list
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'])
s
```

Out[2]:
```
0                 7
1       Heisenberg
2             3.14
3      -1789710578
4    Happy Eating!
dtype: object
```

In [4]:
```
d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900, 'San Francisco': 1100,
     'Austin': 450, 'Boston': None}
cities = pd.Series(d)
cities
```

Out[4]:
```
Austin            450
Boston            NaN
Chicago          1000
New York         1300
Portland          900
San Francisco    1100
dtype: float64
```

Series slicing

```
In [15]: cities[2:4]

Out[15]: city
         Chicago      1000
         New York     1300
```

```
In [16]: cities['Chicago']

Out[16]: 1000.0
```

```
In [17]: cities[cities > 1000]

Out[17]: city
         New York          1300
         San Francisco     1100
```

```
In [20]: cities[cities.isnull()]

Out[20]: city
         Boston   NaN
         dtype: float64
```

- Standard Python slicing
- Using the index
- Boolean slicing

Series operations

- Changing values
- Math

```
In [19]: cities[cities < 1000] = 750
         cities[cities == 750]

Out[19]: city
         Austin       750
         Portland     750
```

```
In [28]: cities[3:] + 7

Out[28]: city
         New York           1307
         Portland            757
         San Francisco      1107
```

```
In [23]: cities[:3] + cities[2:]

Out[23]: city
         Austin              NaN
         Boston              NaN
         Chicago            2000
         New York            NaN
         Portland            NaN
         San Francisco       NaN
         dtype: float64
```

DataFrames

In [72]:
```
d = [{'losses': 5, 'year': 2010, 'wins': 11, 'team': 'Bears'},
     {'losses': 8, 'year': 2011, 'wins': 8, 'team': 'Bears'},
     {'losses': 6, 'year': 2012, 'wins': 10, 'team': 'Bears'},
     {'losses': 1, 'year': 2011, 'wins': 15, 'team': 'Packers'},
     {'losses': 5, 'year': 2012, 'wins': 11, 'team': 'Packers'},
     {'losses': 10, 'year': 2010, 'wins': 6, 'team': 'Lions'},
     {'losses': 6, 'year': 2011, 'wins': 10, 'team': 'Lions'},
     {'losses': 12, 'year': 2012, 'wins': 4, 'team': 'Lions'}]
teams = pd.DataFrame(d, columns=['year', 'team', 'wins', 'losses'])
teams
```

Out[72]:

	year	team	wins	losses
0	2010	Bears	11	5
1	2011	Bears	8	8
2	2012	Bears	10	6
3	2011	Packers	15	1
4	2012	Packers	11	5
5	2010	Lions	6	10
6	2011	Lions	10	6
7	2012	Lions	4	12

Indexes

They’re not columns.

Demo Time!

(Live coding is such a bad idea …)